

# The Numpy array object

## Section contents

- What are Numpy and Numpy arrays?
- Reference documentation
- Import conventions
- Creating arrays
- Functions for creating arrays
- Basic data types
- Basic visualization
- Indexing and slicing
- Copies and views
- Fancy indexing

## What are Numpy and Numpy arrays?

### Python objects

- high-level number objects: integers, floating point
- containers: lists (costless insertion and append), dictionaries (fast lookup)

### Numpy provides

- extension package to Python for multi-dimensional arrays
- closer to hardware (efficiency)
- designed for scientific computation (convenience)
- Also known as *array oriented computing*

```
In [1]: import numpy as np
        a = np.array([0, 1, 2, 3])
        a
```

```
Out[1]: array([0, 1, 2, 3])
```

For example, An array containing:

- values of an experiment/simulation at discrete time steps
- signal recorded by a measurement device, e.g. sound wave
- pixels of an image, grey-level or colour
- 3-D data measured at different X-Y-Z positions, e.g. MRI scan
- ...

**Why it is useful:** Memory-efficient container that provides fast numerical operations.

```
In [2]: L = range(1000)
```

```
In [3]: %timeit [i**2 for i in L]
```

10000 loops, best of 3: 71.6 µs per loop

```
In [4]: a = np.arange(1000)
```

```
In [5]: %timeit a**2
```

100000 loops, best of 3: 1.78 µs per loop

## Reference documentation

- On the web: <http://docs.scipy.org> (<http://docs.scipy.org/>)
- Interactive help:

In [6]: `np.array?`

- Looking for something:

In [7]: `np.lookfor('create array')`

Search results for 'create array'

-----  
`numpy.array`

Create an array.

`numpy.memmap`

Create a memory-map to an array stored in a \*binary\* file on disk.

`numpy.diagflat`

Create a two-dimensional array with the flattened input as a diagonal.

`numpy.fromiter`

Create a new 1-dimensional array from an iterable object.

`numpy.partition`

Return a partitioned copy of an array.

`numpy.ma.diagflat`

Create a two-dimensional array with the flattened input as a diagonal.

`numpy.ctypeslib.as_array`

Create a numpy array from a ctypes array or a ctypes POINTER.

`numpy.ma.make_mask`

Create a boolean mask from an array.

`numpy.ctypeslib.as_ctypes`

Create and return a ctypes object from a numpy array. Actually

`numpy.ma.mrecords.fromarrays`

Creates a mrecarray from a (flat) list of masked arrays.

`numpy.lib.format.open_memmap`

Open a .npy file as a memory-mapped array.

`numpy.ma.MaskedArray.__new__`

Create a new masked array from scratch.

`numpy.lib.arrayterator.Arrayterator`

Buffered iterator for big arrays.

`numpy.ma.mrecords.fromtextfile`

Creates a mrecarray from data stored in the file `filename`.

`numpy.oldnumeric.ma.fromfunction`

apply f to s to create array as in umath.

`numpy.oldnumeric.ma.masked_object`

Create array masked where exactly data equal to value

`numpy.oldnumeric.ma.masked_values`

Create a masked array; mask is nomask if possible.

`numpy.asarray`

Convert the input to an array.

`numpy.ndarray`

`ndarray(shape, dtype=float, buffer=None, offset=0,`

`numpy.recarray`

Construct an ndarray that allows field access using attributes.

`numpy.chararray`

`chararray(shape, itemsize=1, unicode=False, buffer=None, offset=0,`

`numpy.pad`

Pads an array.

`numpy.sum`

Sum of array elements over a given axis.

`numpy.asanyarray`

Convert the input to an ndarray, but pass ndarray subclasses through.

`numpy.copy`  
Return an array copy of the given object.

`numpy.diag`  
Extract a diagonal or construct a diagonal array.

`numpy.load`  
Load an array(s) or pickled objects from `.npy`, `.npz`, or pickled files.

`numpy.sort`  
Return a sorted copy of an array.

`numpy.array_equiv`  
Returns True if input arrays are shape consistent and all elements equal.

`numpy.dtype`  
Create a data type object.

`numpy.choose`  
Construct an array from an index array and a set of arrays to choose from.

`numpy.nditer`  
Efficient multi-dimensional iterator object to iterate over arrays.

`numpy.swapaxes`  
Interchange two axes of an array.

`numpy.full_like`  
Return a full array with the same shape and type as a given array.

`numpy.ones_like`  
Return an array of ones with the same shape and type as a given array.

`numpy.empty_like`  
Return a new array with the same shape and type as a given array.

`numpy.zeros_like`  
Return an array of zeros with the same shape and type as a given array.

`numpy.asarray_chkfinite`  
Convert the input to an array, checking for NaNs or Infs.

`numpy.diag_indices`  
Return the indices to access the main diagonal of an array.

`numpy.ma.choose`  
Use an index array to construct a new array from a set of choices.

`numpy.chararray.tolist`  
`a.tolist()`

`numpy.matlib.rand`  
Return a matrix of random values with given shape.

`numpy.savez_compressed`  
Save several arrays into a single file in compressed ``.npz`` format.

`numpy.ma.empty_like`  
Return a new array with the same shape and type as a given array.

`numpy.ma.make_mask_none`  
Return a boolean mask of the given shape, filled with False.

`numpy.ma.mrecords.fromrecords`  
Creates a MaskedRecords from a list of records.

`numpy.around`  
Evenly round to the given number of decimals.

`numpy.source`  
Print or write to a file the source code for a Numpy object.

`numpy.diagonal`  
Return specified diagonals.

`numpy.histogram2d`  
Compute the bi-dimensional histogram of two data samples.

`numpy.fft.ifft`  
Compute the one-dimensional inverse discrete Fourier Transform.

`numpy.fft.ifftn`  
Compute the N-dimensional inverse discrete Fourier Transform.

`numpy.busdaycalendar`  
A business day calendar object that efficiently stores information

In [10]: `np.con*`

## Import conventions

The general convention to import numpy is:

```
In [11]: import numpy as np
```

Using this style of import is recommended.

## Creating arrays

- **1-D:**

```
In [12]: a = np.array([0, 1, 2, 3])  
a
```

```
Out[12]: array([0, 1, 2, 3])
```

```
In [13]: a.ndim
```

```
Out[13]: 1
```

```
In [14]: a.shape
```

```
Out[14]: (4,)
```

```
In [15]: len(a)
```

```
Out[15]: 4
```

- **2-D, 3-D, ...:**

```
In [16]: b = np.array([[0, 1, 2], [3, 4, 5]])  # 2 x 3 array  
b
```

```
Out[16]: array([[0, 1, 2],  
               [3, 4, 5]])
```

```
In [17]: b.ndim
```

```
Out[17]: 2
```

```
In [18]: b.shape
```

```
Out[18]: (2, 3)
```

```
In [19]: len(b)  # returns the size of the first dimension
```

```
Out[19]: 2
```

```
In [20]: c = np.array([[[1], [2]], [[3], [4]]])  
c
```

```
Out[20]: array([[[1],  
                [2]],  
               [[3],  
                [4]]])
```

```
In [21]: c.shape
```

```
Out[21]: (2, 2, 1)
```

## Exercise: Simple arrays

- Create simple one and two dimensional arrays. First, redo the examples from above. And then create your own.
- Use the functions len, shape and ndim on some of those arrays and observe their output.

## Functions for creating arrays

In practice, we rarely enter items one by one...

- Evenly spaced:

```
In [22]: a = np.arange(10) # 0 .. n-1 (!)
a
```

```
Out[22]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [23]: b = np.arange(1, 9, 2) # start, end (exclusive), step
b
```

```
Out[23]: array([1, 3, 5, 7])
```

- or by number of points:

```
In [24]: c = np.linspace(0, 1, 6) # start, end, num-points
c
```

```
Out[24]: array([ 0. ,  0.2,  0.4,  0.6,  0.8,  1. ])
```

```
In [25]: d = np.linspace(0, 1, 5, endpoint=False)
d
```

```
Out[25]: array([ 0. ,  0.2,  0.4,  0.6,  0.8])
```

- Common arrays:

```
In [26]: a = np.ones((3, 3)) # reminder: (3, 3) is a tuple
a
```

```
Out[26]: array([[ 1.,  1.,  1.],
               [ 1.,  1.,  1.],
               [ 1.,  1.,  1.]])
```

```
In [27]: b = np.zeros((2, 2))
b
```

```
Out[27]: array([[ 0.,  0.],
               [ 0.,  0.]])
```

```
In [28]: c = np.eye(3)
c
```

```
Out[28]: array([[ 1.,  0.,  0.],
               [ 0.,  1.,  0.],
               [ 0.,  0.,  1.]])
```

```
In [29]: d = np.diag(np.array([1, 2, 3, 4]))
d
```

```
Out[29]: array([[1, 0, 0, 0],
```

```
Out[29]: array([[1, 0, 0, 0],
               [0, 2, 0, 0],
               [0, 0, 3, 0],
               [0, 0, 0, 4]])
```

- np.random random numbers (Mersenne Twister PRNG):

```
In [30]: a = np.random.rand(4)    # uniform in [0, 1]
a
```

```
Out[30]: array([ 0.90103456,  0.04550118,  0.26387871,  0.42707105])
```

```
In [31]: b = np.random.randn(4)   # Gaussian
b
```

```
Out[31]: array([-0.24637953,  0.5029874 ,  1.18003342,  0.06942228])
```

```
In [32]: np.random.seed(1234)     # Setting the random seed
```

## Exercise: Creating arrays using functions

- Experiment with arange, linspace, ones, zeros, eye and diag.
- Create different kinds of arrays with random numbers.
- Try setting the seed before creating an array with random values.
- Look at the function np.empty. What does it do? When might this be useful?

## Basic data types

You may have noticed that, in some instances, array elements are displayed with a trailing dot (e.g. 2. vs 2). This is due to a difference in the data-type used:

```
In [33]: a = np.array([1, 2, 3])
a.dtype
```

```
Out[33]: dtype('int64')
```

```
In [34]: b = np.array([1., 2., 3.])
b.dtype
```

```
Out[34]: dtype('float64')
```

## Tip

Different data-types allow us to store data more compactly in memory, but most of the time we simply work with floating point numbers. Note that, in the example above, NumPy auto-detects the data-type from the input.

You can explicitly specify which data-type you want:

```
In [35]: c = np.array([1, 2, 3], dtype=float)
c.dtype
```

```
Out[35]: dtype('float64')
```

The **default** data type is floating point:

```
In [36]: a = np.ones((3, 3))
a.dtype
```

```
Out[36]: dtype('float64')
```

There are also other types:

Complex

```
In [37]: d = np.array([1+2j, 3+4j, 5+6*1j])
         d.dtype
```

```
Out[37]: dtype('complex128')
```

Bool

```
In [38]: e = np.array([True, False, False, True])
         e.dtype
```

```
Out[38]: dtype('bool')
```

Strings

```
In [39]: f = np.array(['Bonjour', 'Hello', 'Halo',])
         f.dtype    # <--- strings containing max. 7 letters
```

```
Out[39]: dtype('S7')
```

Much more

- int32
- int64
- unit32
- unit64

## Basic visualization

### Tip

Now that we have our first data arrays, we are going to visualize them.

Start by launching IPython in *pylab* mode.

```
$ ipython --pylab
```

Or the notebook:

```
$ ipython notebook --pylab=inline
```

Alternatively, if IPython has already been started:

```
In [40]: %pylab

Using matplotlib backend: TkAgg
Populating the interactive namespace from numpy and matplotlib

WARNING: pylab import has clobbered these variables: ['e', 'f']
`%pylab --no-import-all` prevents importing * from pylab and numpy
```

Or, from the notebook:

```
In []: %pylab inline
```

The inline is important for the notebook, so that plots are displayed in the notebook and not in a new window.

*Matplotlib* is a 2D plotting package. We can import its functions as below:

```
In [41]: import matplotlib.pyplot as plt # the tidy way
```

And then use (note that you have to use `show` explicitly):

```
In [42]: plt.plot(x, y)    # line plot
```

```
plt.show()      # <-- shows the plot (not needed with pylab)
```

```
-----  
NameError                                Traceback (most recent call last)  
<ipython-input-42-cca45a0ba107> in <module>()  
----> 1 plt.plot(x, y)      # line plot  
      2 plt.show()         # <-- shows the plot (not needed with pylab)  
  
NameError: name 'x' is not defined
```

Or, if you are using *pylab*:

```
In []: plot(x, y)      # line plot
```

Using `import matplotlib.pyplot as plt` is recommended for use in scripts. Whereas *pylab* is recommended for interactive exploratory work.

- **1D plotting:**

```
In []: x = np.linspace(0, 3, 20)  
      y = np.linspace(0, 9, 20)  
      plt.plot(x, y)      # line plot
```

```
In []: plt.plot(x, y, 'o') # dot plot
```

- **2D arrays** (such as images):

```
In []: image = np.random.rand(30, 30)  
      plt.imshow(image, cmap=plt.cm.hot)  
      plt.colorbar()
```

More in the Matplotlib tutorial this afternoon

## Exercise: Simple visualizations

- Plot some simple arrays.
- Try to use both the IPython shell and the notebook, if possible.
- Try using the gray colormap.

## Indexing and slicing

The items of an array can be accessed and assigned to the same way as other Python sequences (e.g. lists):

```
In [43]: a = np.arange(10)  
      a
```

```
Out[43]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [44]: a[0], a[2], a[-1]
```

```
Out[44]: (0, 2, 9)
```

## Warning

Indices begin at 0, like other Python sequences (and C/C++). In contrast, in Fortran or Matlab, indices begin at 1.

The usual python idiom for reversing a sequence is supported:

```
In [45]: a[::-1]
```



```
Out[45]: array([9, 8, 7, 6, 5, 4, 3, 2, 1, 0])
```

For multidimensional arrays, indexes are tuples of integers:

```
In [46]: a = np.diag(np.arange(3))  
a
```

```
Out[46]: array([[0, 0, 0],  
               [0, 1, 0],  
               [0, 0, 2]])
```

```
In [47]: a[1, 1]
```

```
Out[47]: 1
```

```
In [48]: a[2, 1] = 10 # third line, second column  
a
```

```
Out[48]: array([[ 0,  0,  0],  
               [ 0,  1,  0],  
               [ 0, 10,  2]])
```

```
In [49]: a[1]
```

```
Out[49]: array([0, 1, 0])
```

Note that:

- In 2D, the first dimension corresponds to rows, the second to columns.
- Let us repeat together: the first dimension corresponds to **rows**, the second to **columns**.
- for multidimensional a, a[0] is interpreted by taking all elements in the unspecified dimensions.

**Slicing** Arrays, like other Python sequences can also be sliced:

```
In [50]: a = np.arange(10)  
a
```

```
Out[50]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [51]: a[2:9:3] # [start:end:step]
```

```
Out[51]: array([2, 5, 8])
```

Note that the last index is not included! :

```
In [52]: a[:4]
```

```
Out[52]: array([0, 1, 2, 3])
```

All three slice components are not required: by default, `start` is 0, `end` is the last and `step` is 1:

```
In [53]: a[1:3]
```

```
Out[53]: array([1, 2])
```

```
In [54]: a[::2]
```

```
Out[54]: array([0, 2, 4, 6, 8])
```

```
In [55]: a[3:]
```

```
Out[55]: array([3, 4, 5, 6, 7, 8, 9])
```

A small illustrated summary of Numpy indexing and slicing...

```
In [56]: from IPython.display import Image  
Image(filename='images/numpy_indexing.png')
```

```

-----
IOError                                Traceback (most recent call last)
<ipython-input-56-ef00be976d20> in <module>()
      1 from IPython.display import Image
----> 2 Image(filename='images/numpy_indexing.png')

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in __init__(self, data, url, filename, format, embed, width, height, retina)
    599     self.height = height
    600     self.retina = retina
--> 601     super(Image, self).__init__(data=data, url=url, filename=filename)
    602
    603     if retina:

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in __init__(self, data, url, filename)
    303     self.filename = None if filename is None else unicode(filename)
    304
--> 305     self.reload()
    306
    307     def reload(self):

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in reload(self)
    621     """Reload the raw data from file or URL."""
    622     if self.embed:
--> 623         super(Image, self).reload()
    624     if self.retina:
    625         self._retina_shape()

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in reload(self)
    308     """Reload the raw data from file or URL."""
    309     if self.filename is not None:
--> 310         with open(self.filename, self._read_flags) as f:
    311             self.data = f.read()
    312     elif self.url is not None:

IOError: [Errno 2] No such file or directory: u'images/numpy_indexing.png'

```

You can also combine assignment and slicing:

```

In [57]: a = np.arange(10)
        a[5:] = 10
        a

```

```

Out[57]: array([ 0,  1,  2,  3,  4, 10, 10, 10, 10, 10])

```

```

In [58]: b = np.arange(5)
        a[5:] = b[::-1]
        a

```

```

Out[58]: array([0, 1, 2, 3, 4, 4, 3, 2, 1, 0])

```

## Exercise: Indexing and slicing

- Try the different flavours of slicing, using start, end and step.
- Verify that the slices in the diagram above are indeed correct. You may use the following expression to create the array:

```

In [59]: np.arange(6) + np.arange(0, 51, 10)[: , np.newaxis]

```

```

Out[59]: array([[ 0,  1,  2,  3,  4,  5],
               [10, 11, 12, 13, 14, 15],
               [20, 21, 22, 23, 24, 25],
               [30, 31, 32, 33, 34, 35],
               [40, 41, 42, 43, 44, 45],
               [50, 51, 52, 53, 54, 55]])

```

- Try assigning a smaller 2D array to a larger 2D array, like in the 1D example above.
- Use a different step, e.g. -2, in the reversal idiom above. What effect does this have?

## Exercise: Array creation

Create the following arrays (with correct data types):

```
[[1, 1, 1, 1], [1, 1, 1, 1], [1, 1, 1, 2], [1, 6, 1, 1]] [[0., 0., 0., 0., 0.], [2., 0., 0., 0., 0.], [0., 3., 0., 0., 0.], [0., 0., 4., 0., 0.], [0., 0., 0., 5., 0.], [0., 0., 0., 0., 6.]]
```

Par on course: 3 statements for each

*Hint:* Individual array elements can be accessed similarly to a list, e.g. `a[1]` or `a[1, 2]`.

*Hint:* Examine the docstring for `diag`.

## Exercise: Tiling for array creation

Skim through the documentation for `np.tile`, and use this function to construct the array:

```
[[4, 3, 4, 3, 4, 3], [2, 1, 2, 1, 2, 1], [4, 3, 4, 3, 4, 3], [2, 1, 2, 1, 2, 1]]
```

## Copies and views

A slicing operation creates a **view** on the original array, which is just a way of accessing array data. Thus the original array is not copied in memory. You can use `np.may_share_memory()` to check if two arrays share the same memory block. Note however, that this uses heuristics and may give you false positives.

**When modifying the view, the original array is modified as well:**

```
In [60]: a = np.arange(10)
a
```

```
Out[60]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [61]: b = a[::2]
b
```

```
Out[61]: array([0, 2, 4, 6, 8])
```

```
In [62]: np.may_share_memory(a, b)
```

```
Out[62]: True
```

```
In [63]: b[0] = 12
b
```

```
Out[63]: array([12, 2, 4, 6, 8])
```

```
In [64]: a # (!)
```

```
Out[64]: array([12, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [65]: a = np.arange(10)
c = a[::2].copy() # force a copy
c[0] = 12
a
```

```
Out[65]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [66]: np.may_share_memory(a, c)
```

```
Out[66]: False
```

This behavior can be surprising at first sight... but it allows to save both memory and time.

## Worked example: Prime number sieve

```
In [67]: from IPython.display import Image
Image(filename='images/prime-sieve.png')

-----
IOError                                Traceback (most recent call last)
<ipython-input-67-08ed76b99731> in <module>()
      1 from IPython.display import Image
----> 2 Image(filename='images/prime-sieve.png')

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in __init__(self, data, url, filename, format, embed, width, height, retina)
    599     self.height = height
    600     self.retina = retina
--> 601     super(Image, self).__init__(data=data, url=url, filename=filename)
    602
    603     if retina:

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in __init__(self, data, url, filename)
    303     self.filename = None if filename is None else unicode(filename)
    304
--> 305     self.reload()
    306
    307     def reload(self):

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in reload(self)
    621     """Reload the raw data from file or URL."""
    622     if self.embed:
--> 623         super(Image, self).reload()
    624     if self.retina:
    625         self._retina_shape()

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in reload(self)
    308     """Reload the raw data from file or URL."""
    309     if self.filename is not None:
--> 310         with open(self.filename, self._read_flags) as f:
    311             self.data = f.read()
    312     elif self.url is not None:

IOError: [Errno 2] No such file or directory: u'images/prime-sieve.png'
```

Compute prime numbers in 0--99, with a sieve

- Construct a shape (100,) boolean array `is_prime`, filled with True in the beginning:

```
In []: is_prime = np.ones((100,), dtype=bool)
```

- Cross out 0 and 1 which are not primes:

```
In []: is_prime[:2] = 0
```

- For each integer `j` starting from 2, cross out its higher multiples:

```
In []: N_max = int(np.sqrt(len(is_prime)))
      for j in range(2, N_max):
          is_prime[2*j::j] = False
```

- Skim through `help(np.nonzero)`, and print the prime numbers
- Follow-up:
  - Move the above code into a script file named `prime_sieve.py`
  - Run it to check it works
  - Use the optimization suggested in [the sieve of Eratosthenes \(http://en.wikipedia.org/wiki/Sieve\\_of\\_Eratosthenes\)](http://en.wikipedia.org/wiki/Sieve_of_Eratosthenes):
  - Skip `j` which are already known to not be primes
  - The first number to cross out is  $j^2$

## Fancy indexing

### Tip

Numpy arrays can be indexed with slices, but also with boolean or integer arrays (**masks**). This method is called *fancy indexing*. It creates **copies not views**.

### Using boolean masks

```
In [68]: np.random.seed(3)
a = np.random.random_integers(0, 20, 15)
a
```

```
Out[68]: array([10,  3,  8,  0, 19, 10, 11,  9, 10,  6,  0, 20, 12,  7, 14])
```

```
In [69]: (a % 3 == 0)
```

```
Out[69]: array([False,  True, False,  True, False, False, False,  True, False,
                True,  True, False,  True, False, False], dtype=bool)
```

```
In [70]: mask = (a % 3 == 0)
extract_from_a = a[mask] # or, a[a%3==0]
extract_from_a      # extract a sub-array with the mask
```

```
Out[70]: array([ 3,  0,  9,  6,  0, 12])
```

Indexing with a mask can be very useful to assign a new value to a sub-array:

```
In [71]: a[a % 3 == 0] = -1
a
```

```
Out[71]: array([10, -1,  8, -1, 19, 10, 11, -1, 10, -1, -1, 20, -1,  7, 14])
```

### Indexing with an array of integers

```
In [72]: a = np.arange(0, 100, 10)
a
```

```
Out[72]: array([ 0, 10, 20, 30, 40, 50, 60, 70, 80, 90])
```

Indexing can be done with an array of integers, where the same index is repeated several time:

```
In [73]: a[[2, 3, 2, 4, 2]] # note: [2, 3, 2, 4, 2] is a Python list
```

```
Out[73]: array([20, 30, 20, 40, 20])
```

New values can be assigned with this kind of indexing:

```
In [74]: a[[9, 7]] = -100
a
```

```
Out[74]: array([ 0, 10, 20, 30, 40, 50, 60, -100, 80, -100])
```

## Tip

When a new array is created by indexing with an array of integers, the new array has the same shape than the array of integers:

```
In [75]: a = np.arange(10)
idx = np.array([[3, 4], [9, 7]])
idx.shape
```

```
Out[75]: (2, 2)
```

```
In [76]: a[idx]
```

```
Out[76]: array([[3, 4],
               [9, 7]])
```

The image below illustrates various fancy indexing applications

```
In [77]: from IPython.display import Image
Image(filename='images/numpy_fancy_indexing.png')
```

```
-----
IOError                                Traceback (most recent call last)
<ipython-input-77-beb616f120d9> in <module>()
      1 from IPython.display import Image
----> 2 Image(filename='images/numpy_fancy_indexing.png')

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in __init__(self, data, url, filename, format, embed, width, height, retina)
    599     self.height = height
    600     self.retina = retina
--> 601     super(Image, self).__init__(data=data, url=url, filename=filename)
    602
    603     if retina:

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in __init__(self, data, url, filename)
    303     self.filename = None if filename is None else unicode(filename)
    304
--> 305     self.reload()
    306
    307     def reload(self):

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in reload(self)
    621     """Reload the raw data from file or URL."""
    622     if self.embed:
--> 623         super(Image, self).reload()
    624     if self.retina:
    625         self._retina_shape()

/usr/lib/python2.7/dist-packages/IPython/core/display.pyc in reload(self)
    308     """Reload the raw data from file or URL."""
    309     if self.filename is not None:
--> 310         with open(self.filename, self._read_flags) as f:
    311             self.data = f.read()
    312     elif self.url is not None:

IOError: [Errno 2] No such file or directory: u'images/numpy_fancy_indexing.png'
```

## Exercise: Fancy indexing

- Again, verify the fancy indexing shown in the diagram above.

- Use fancy indexing on the left and array creation on the right to assign values from a smaller array to a larger array.